



SQL Tuning Tips You Can't Do Without



Maria Colgan

Distinguished Product Manager



@SQLMaria

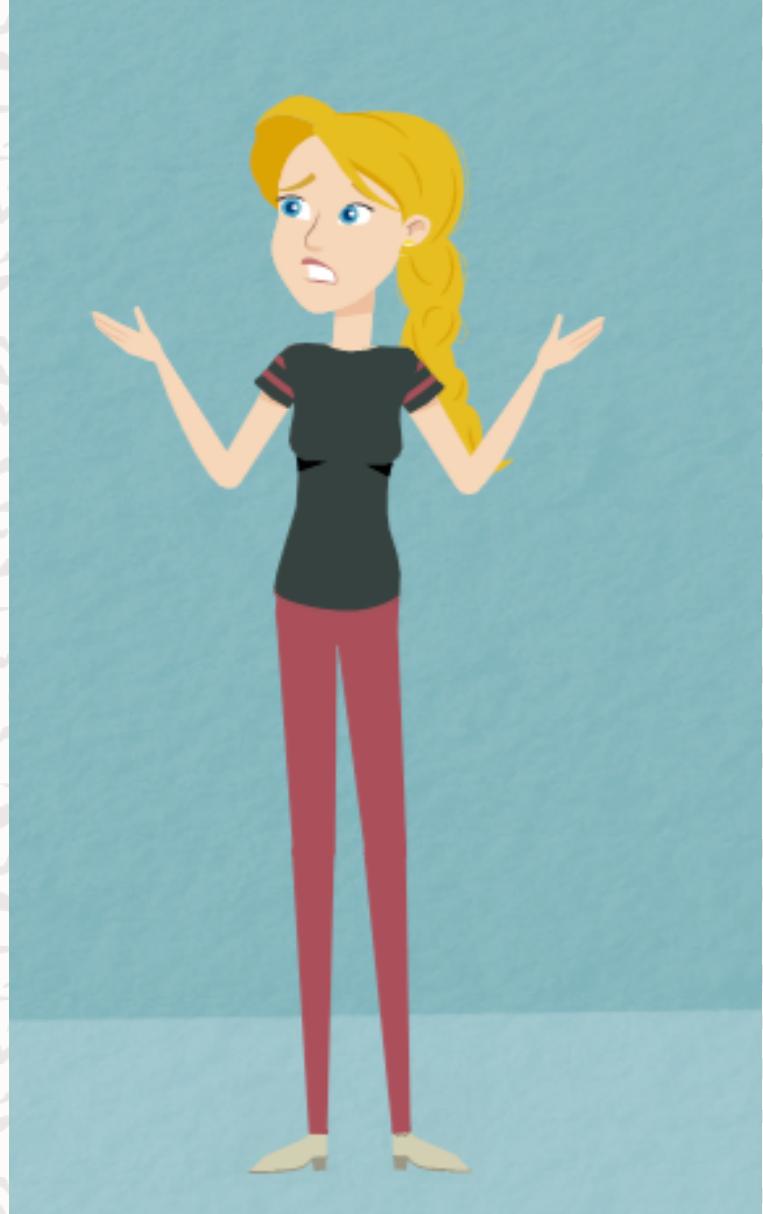


Disclaimer

The goal of this session to provide you tips on how to correct common problems identified from a SQL execution plans

This session will not provide you with sudden enlightenment, making you an Optimizer expert or give you the power to tune SQL statements with the flick of your wrist!

Problem Statement 1



What Join Method Should The Optimizer Pick?

```
SELECT      p.prod_name, SUM(s.quantity_sold)  
FROM        sales s, products p  
WHERE       s.prod_id = p.prod_id  
AND         s.time_id = '03-DEC-20';  
GROUP BY    p.prod_name;
```

SALES table has 10 million rows

PRODUCTS table has 100,000 rows

Got a Nested Loops Join Instead of Hash Join

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-DEC-20';
GROUP BY    p.prod_name;
```

SALES table has 10 million rows

PRODUCTS table has 100,000 rows

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	TABLE ACCESS STORAGE FULL	SALES	0	1
*	INDEX UNIQUE SCAN	PRODUCTS_PK	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

Got a Nested Loops Join Instead of Hash Join

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-DEC-20';
GROUP BY    p.prod_name;
```

1. Always start by looking at the Cardinality Estimate

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	5	TABLE ACCESS STORAGE FULL	0	1
*	6	INDEX UNIQUE SCAN	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

Got a Nested Loops Join Instead of Hash Join

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-DEC-20';
GROUP BY    p.prod_name;
```

1. Always start by looking at the Cardinality Estimate
2. The zero's in the Starts column also look suspicious

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	TABLE ACCESS STORAGE FULL	SALES	0	1
*	INDEX UNIQUE SCAN	PRODUCTS_PK	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

Need to Gather More Information

```
SELECT table_name, partitioned  
FROM   user_tables  
WHERE  table_name in  
       ( 'SALES', 'PRODUCTS' );
```

TABLE_NAME	PARTITIONED
PRODUCTS	NO
SALES	YES

Need to Gather More Information

```
SELECT table_name, partitioned  
FROM   user_tables  
WHERE  table_name in  
       ('SALES','PRODUCTS');
```

TABLE_NAME	PARTITIONED
PRODUCTS	NO
SALES	YES

```
SELECT table_name, stale_stats  
FROM   user_tab_statistics  
WHERE  table_name in ('SALES','PRODUCTS');
```

TABLE_NAME	STALE_STATS
PRODUCTS	NO
SALES	NO
SALES	

Need to Gather More Information

```
SELECT table_name, partitioned
FROM   user_tables
WHERE  table_name in
      ('SALES','PRODUCTS');
```

TABLE_NAME	PARTITIONED
PRODUCTS	NO
SALES	YES

```
SELECT table_name, stale_stats
FROM   user_tab_statistics
WHERE  table_name in ('SALES','PRODUCTS');
```

TABLE_NAME	STALE_STATS
PRODUCTS	NO

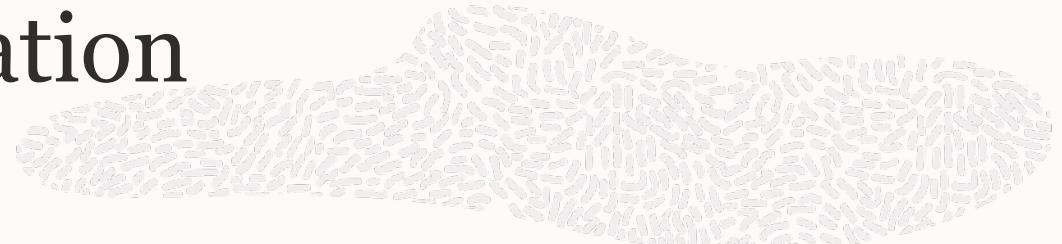
SALES

NO — Table Stats

Partition Level Stats

Need to Gather More Information

Next Look at the Where Clause Predicates



AND s.time_id = '03-DEC-20' ;

```
SELECT min(time_id),  
       max(time_id)  
FROM   sales;
```

MIN(DATE_ID) MAX(DATE_ID)

01-JAN-04 03-DEC-20

Need to Gather More Information

Next Look at the Where Clause Predicates

AND s.time_id = '03-DEC-20' ;

```
SELECT min(time_id),  
       max(time_id)  
FROM   sales;
```

MIN(DATE_ID) **MAX(DATE_ID)**

01-JAN-04 03-DEC-20

```
DECLARE  
  rv RAW(32) ;  
  dt DATE;  
BEGIN  
  SELECT high_value INTO rv  
  FROM   user_tab_col_statistics  
  WHERE  table_name='SALES'  
  AND    column_name='TIME_ID';  
  dbms_stats.convert_raw_value(rv, dt);  
  dbms_output.put_line( TO_CHAR(dt,'dd-MON-yy'));  
END;  
/
```

31-OCT-2020

PL/SQL procedure successfully completed.



Stale Statistics Cause an “Out of Range” Issues

- Statistics for `TIME_ID` column show max value as **‘31-OCT-2020’**
- Our query is looking for sales on **‘03-DEC-2020’**
- Optimizer checks if predicate falls between the min, max value of column

Stale Statistics Cause an “Out of Range” Issues

- Statistics for `TIME_ID` column show max value as **‘31-OCT-2020’**
- Our query is looking for sales on **‘03-DEC-2020’**
- Optimizer checks if predicate falls between the min, max value of column
- “Out of Range” means value supplied is outside the min, max range
- Optimizer prorates cardinality based on the distance between the predicate value and the maximum value

Stale Statistics Cause an “Out of Range” Issues

- Statistics for `TIME_ID` column show max value as **‘31-OCT-2020’**
- Our query is looking for sales on **‘03-DEC-2020’**
- Optimizer checks if predicate falls between the min, max value of column
- “Out of Range” means value supplied is outside the min, max range
- Optimizer prorates cardinality based on the distance between the predicate value and the maximum value

In **10053 trace file** you will find:

“Using prorated density: 0.000000 of col #1 as selectivity of out-of-range/non-existent value pred”

Stale Statistics Cause an “Out of Range” Issues

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-DEC-20';
GROUP BY    p.prod_name;
```

Optimizer assumes there are no rows for time_id='03-DEC-20' because it is out side of the [MIN,MAX] range in the statistics

Id	Operation	Name	Starts	Rows
0	SELECT STATEMENT		1	1
1	HASH GROUP BY		1	1
2	NESTED LOOPS		1	1
3	NESTED LOOPS		1	1
4	PARTITION RANGE SINGLE		1	1
*	5	TABLE ACCESS STORAGE FULL	0	1
*	6	INDEX UNIQUE SCAN	0	1
7	TABLE ACCESS BY INDEX ROWID	PRODUCTS	0	1

Always be wary of a cardinality estimates of 1 !

Fix Out-of-Range Cardinality Misestimate



- By default stats are only considered stale after 10% of the rows have changed
- Our tables SALES has 10 million rows in it
 - 1 million rows need to be insert before the statistics are considered stale

Fix Out-of-Range Cardinality Misestimate

Option 1 - USE DBMS_STATS.SET_TABLE_PREFS To Change Staleness Threshold



- By default stats are only considered stale after 10% of the rows have changed
- Our tables SALES has 10 million rows in it
 - 1 million rows need to be insert before the statistics are considered stale
- To avoid out-of-range problems lower the staleness threshold for larger tables

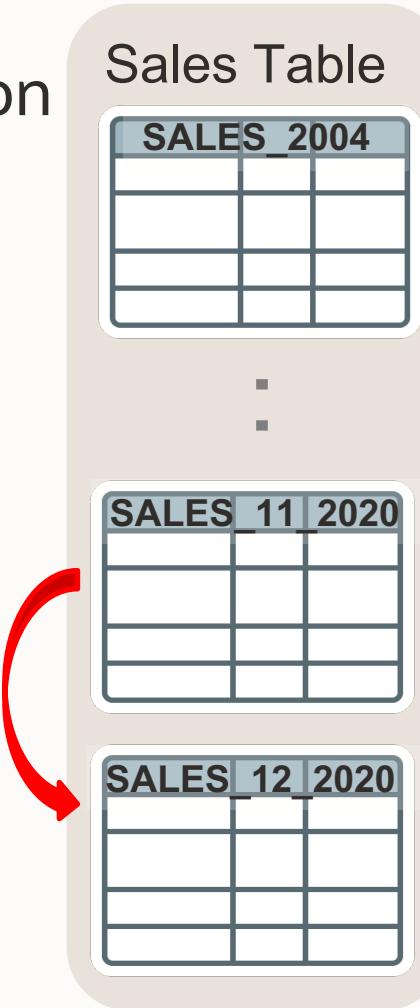
```
BEGIN
    dbms_stats.set_table_prefs('SH', 'SALES', 'STALE_PERCENT', '1');
    dbms_stats.gather_table_stats('SH', 'SALES');
END;
/
```

Fix Out-of-Range Cardinality Misestimate

Option 2 - USE DBMS_STATS.COPY_TABLE_STATS()



- Copies stats from source partition to destination partition
- Adjusts min & max values for partition column at both partition & global level
- Copies statistics of the dependent objects
 - Columns, local indexes etc.
 - Does not update global indexes



```
dbms_stats.copy_table_stats  
( 'SH',  
  'SALES',  
  'SALES_11_2020',  
  'SALES_12_2020' );
```

Fixing Out-of-Range Error Automatic Hash Join Plan

```
SELECT      p.prod_name, SUM(s.quantity_sold)
FROM        sales3 s, products p
WHERE       s.prod_id = p.prod_id
AND         s.time_id = '03-DEC-20';
GROUP BY    p.prod_name;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT		1	32	81 (20)		
1	HASH GROUP BY		1	32	81 (20)		
*	HASH JOIN		3671	114K	80 (19)		
3	JOIN FILTER CREATE	:BF0000	3671	114K	80 (19)		
4	PARTITION RANGE SINGLE		3596	61132	17 (12)	KEY	KEY
*	TABLE ACCESS STORAGE FULL	SALES	3596	61132	17 (12)	KEY	KEY
6	JOIN FILTER USE	:BF0000	1398K	20M	58 (14)		
*	TABLE ACCESS STORAGE FULL	PRODUCTS	1398K	20M	58 (14)		

Problem Statement 2



What Index Do You Think The Optimizer Should Pick?

```
SELECT comment  
FROM my_sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```

MY_SALES table has **2 million rows**

Two Index:

PROD_CUST_COM_IND (prod_id, comment, cust_id)

PROD_CUST (prod_id, cust_id)

Wrong Index is Chosen

```
SELECT comment  
FROM my_sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```

MY_SALES table has 2 million rows

Two Index:

PROD_CUST_COM_IND (prod_id, comment, cust_id)

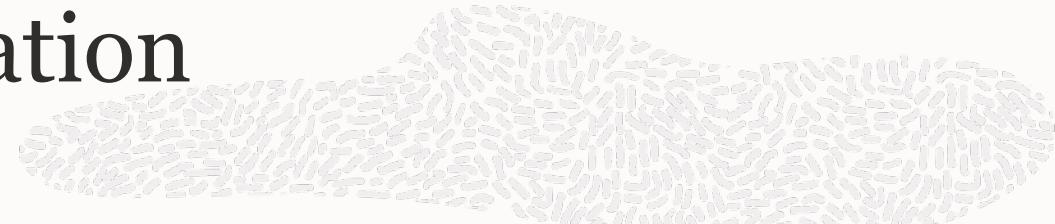
PROD_CUST (prod_id, cust_id)

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	2580 (1)
1	TABLE ACCESS BY INDEX ROWID	MY_SALES	2579	2580 (1)
2	INDEX RANGE SCAN	PROD_CUST_IND	2579	9 (1)

WHY DID IT NOT PICK THE INDEX WITH ALL THE COLUMNS NEEDED?

Need to Gather More Information

Confirm our Preferred Plan is Possible



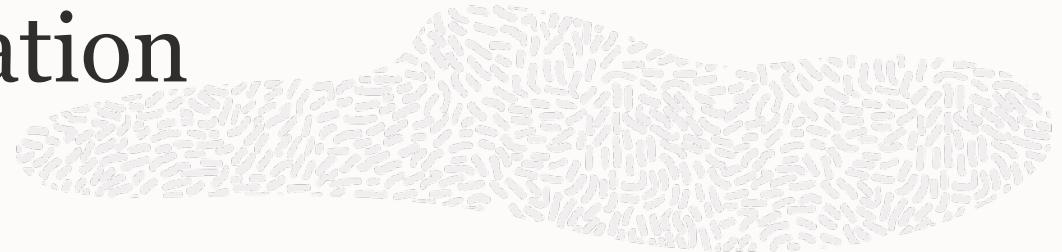
```
SELECT /* INDEX(PROD_CUST_COM_IND) */ comment  
FROM   my_sales  
WHERE  prod_id = 141  
AND    cust_id < 8938;
```

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	9728 (1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728 (1)

WHY IS THE COST SO HIGH?

Need to Gather More Information

Check Table Stats



```
SELECT table_name, partitioned  
FROM user_tables  
WHERE table_name = 'MY_SALES';
```

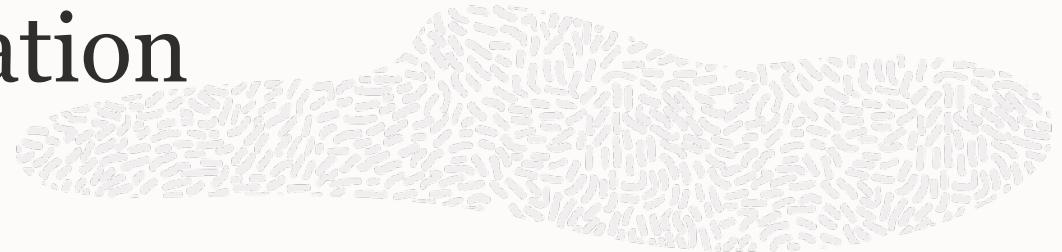
```
SELECT table_name, stale_stats  
FROM user_tab_statistics  
WHERE table_name = 'MY_SALES';
```

TABLE_NAME	PARTITIONED
MY_SALES	NO

TABLE_NAME	STALE_STATS
MY_SALES	NO

Need to Gather More Information

Check Index Stats

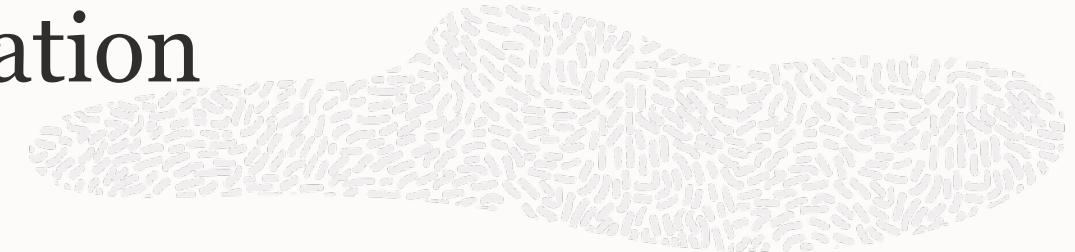


```
SELECT index_name, stale_stats  
FROM   user_ind_statistics  
WHERE  table_name = 'MY_SALES';
```

INDEX_NAME	STALE_STATS
PROD_CUST_COM_IND	NO
PROD_CUST	NO

Need to Gather More Information

Check Index Stats



```
SELECT index_name, stale_stats  
FROM user_ind_statistics  
WHERE table_name = 'MY_SALES';
```

INDEX_NAME	STALE_STATS
PROD_CUST_COM_IND	NO
PROD_CUST	NO

```
SELECT index_name, leaf_blocks, blevel  
FROM user_indexes  
WHERE table_name = 'MY_SALES';
```

INDEX_NAME	LEAF_BLOCKS	BLEVEL
PROD_CUST_COM_IND	699467	9
PROD_CUST	5468	2

Need to Gather More Information

Look At Predicate Information Under the Plan



```
SELECT /* INDEX(PROD_CUST_COM_IND) */ comment  
FROM my_sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	9728 (1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728 (1)

Predicate Information (identified by operation id):

```
-----  
1 - access("PROD_ID"=141 AND "CUST_ID"<8938)  
      filter("CUST_ID"<8938)
```

Remember the order of the column
in the PROD_CUST_COM_IND index
(prod_id, comment, cust_id)

Additional Information Under The Execution Plan

```
SELECT /*+ gather_plan_statistics */ count(*) FROM sales2 WHERE prod_id=to_number('139')
```

Plan hash value: 1631620387

Id	Operation	Name	Starts	E-Rows	Cost (%CPU)	A-Rows
0	SELECT STATEMENT		1	1	35 (100)	1
1	SORT AGGREGATE		1	1		1
* 2	INDEX RANGE SCAN	MY_PROD_IND	1	12762	35 (0)	11574

Predicate Information (identified by operation id):

2 - access("PROD_ID"=139)

Access predicate

- Where clause predicate used for data retrieval
 - The start and stop keys for an index
 - If rowids are passed to a table scan

Additional Information Under The Execution Plan

```
SQL> SELECT username  
2   FROM my_users  
3 WHERE username LIKE 'MAR%';
```

USERNAME

MARIA

Plan hash value: 2982854235

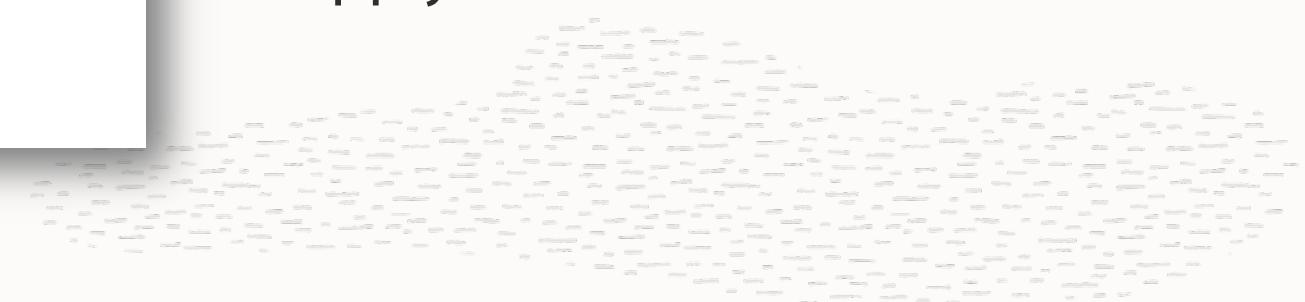
Id Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0 SELECT STATEMENT				2 (100)	
* 1 TABLE ACCESS FULL MY_USERS 1 66 2 (0) 00:00:01					

Predicate Information (identified by operation id):

```
1 - filter("USERNAME" LIKE 'MAR%')
```

Filter predicate

- Where clause predicate that is not used for data retrieval but to eliminate uninteresting row once the data is found
- Requires additional CPU resources to apply filters



Costing of Each Index Access is Different



Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	2580	(1)
1	TABLE ACCESS BY INDEX ROWID	MY_SALES	2579	2580	(1)
2	INDEX RANGE SCAN	PROD_CUST_IND	2579	9	(1)

Predicate Information (identified by operation id):

```
1 - access("PRODUCT_ID"=517538 AND "CUST_ID"<8938)
```

Both columns are used as **Access Predicates**

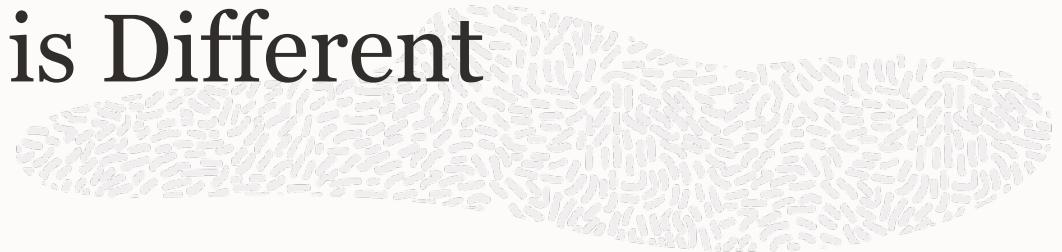
Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	9728	(1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728	(1)

Predicate Information (identified by operation id):

```
1 - access("PRODUCT_ID"=517538 AND "CUST_ID"<8938)
    - filter("CUST_ID"<8938)
```

Only the PROD_ID is used as **Access Predicates**

Costing of Each Index Access is Different



Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	2580 (1)
1	TABLE ACCESS BY INDEX ROWID	MY_SALES	2579	2580 (1)
2	INDEX RANGE SCAN	PROD_CUST_IND	2579	9 (1)

Predicate Information (identified by operation id):

1 - access("PRODUCT_ID"=517538 AND "CUST_ID"<8938)

Both columns are used as **Access Predicates**

Cost calculation for this index:

blevel + (cardinality X leaf blocks)
#rows

$$2 + (\underline{2579} \times 5468) = 9 \\ 2098400$$

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	9728 (1)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND	2579	9728 (1)

Predicate Information (identified by operation id):

1 - access("PRODUCT_ID"=517538 AND "CUST_ID"<8938)
- filter("CUST_ID"<8938)

Only the PROD_ID is used as **Access Predicates**

Cost calculation for this index:

blevel + (1 X Leaf Block)
NDV of col1

$$9 + (0.013 \times 699467) = 9724$$

The CPU cost of applying the filter makes up the additional cost of 4

Common Misconception

Common misconception Optimizer
will pick the index with all the
necessary columns in it



Common Misconception

Common misconception Optimizer will pick the index with all the necessary columns in it

Optimizer picks the index based on the cost of the access



Common Misconception

Common misconception Optimizer will pick the index with all the necessary columns in it

Optimizer picks the index based on the cost of the access

How the WHERE clause predicates are used and the order of the columns in the index have a massive impact on the Optimizers choice



Solution: Right Index is Chosen

```
SELECT comment  
FROM my_sales  
WHERE prod_id = 141  
AND cust_id < 8938;
```

MY_SALES table has 2 million rows

Three Index:

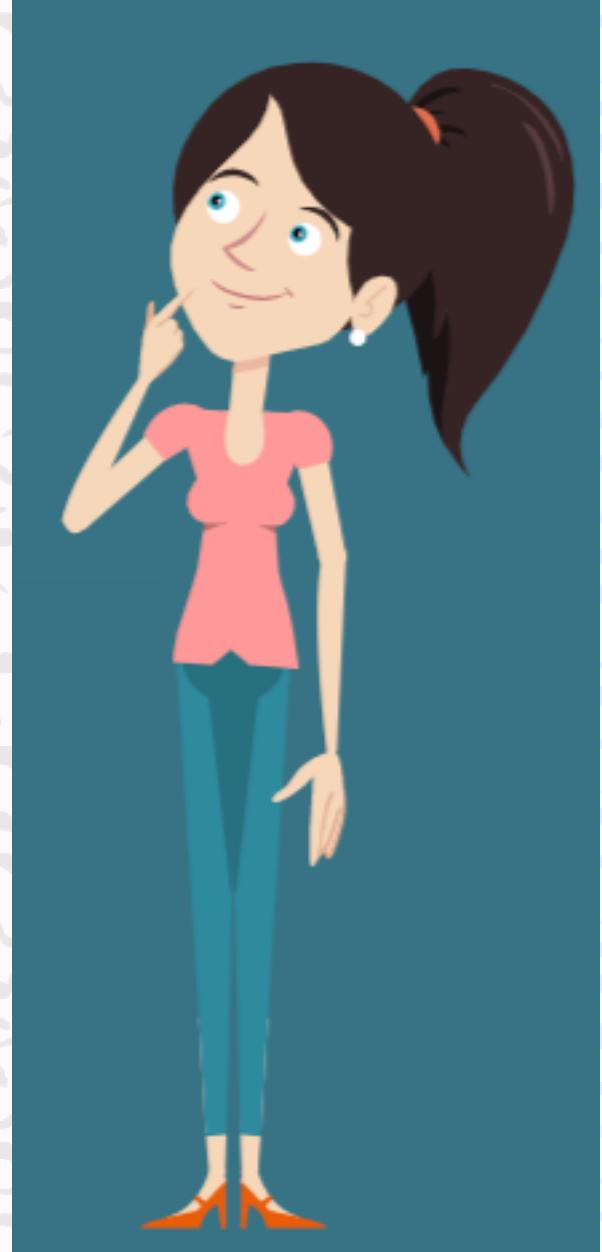
PROD_CUST_COM_IND (prod_id, comment, cust_id)

PROD_CUST (prod_id, cust_id)

PROD_CUST_COM_IND2(prod_id, cust_id, comment)

Id	Operation	Name	Rows	COST (%CPU)
0	SELECT STATEMENT		2579	1935 (100)
1	INDEX RANGE SCAN	PROD_CUST_COM_IND2	2579	1935 (0)

Problem Statement 3



What Join Method Should the Optimizer Pick?

```
SELECT COUNT(*)  
FROM sales s, customers c  
WHERE s.cust_id = c.cust_id  
AND substr(c.cust_state_province,1,2)='CA';
```

SALES table has 10 million rows
CUSTOMERS table has 10,000 rows
Sales table has index on CUST_ID

Got a Nested Loops Join Instead of Hash Join

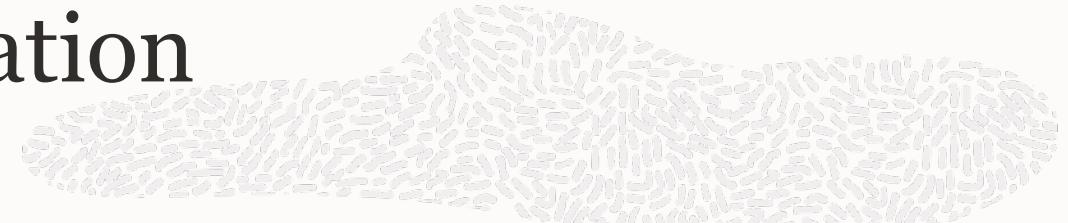
```
SELECT COUNT(*)  
FROM sales s, customers c  
WHERE s.cust_id = c.cust_id  
AND substr(c.cust_state_province,1,2)='CA';
```

SALES table has 10 million rows
CUSTOMERS table has 10,000 rows
Sales table has index on CUST_ID

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			26 (100)
1	SORT AGGREGATE		1	
2	NESTED LOOPS		41	26 (0)
3	TABLE ACCESS FULL	CUSTOMER	100	2 (0)
* 4	BITMAP CONVERSION COUNT		4	26 (0)
* 5	BITMAP INDEX SINGLE VALUE	SALES_CUST_BIX		

Need to Gather More Information

Confirm our Preferred Plan is Possible



```
SELECT /*+ USE_HASH(s) */ COUNT(*)  
FROM      sales s, customers c  
WHERE     s.cust_id = c.cust_id  
AND       substr(c.cust_state_province,1,2)='CA';
```

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			1114 (100)
1	SORT AGGREGATE		1	
2	HASH JOIN		14732	1114 (2)
* 4	TABLE ACCESS FULL	CUSTOMER	100	406 (1)
* 5	TABLE ACCESS FULL	SALES	10M	703 (2)

WHY IS THE COST SO HIGH?

Need to Gather More Information

Compare Estimated Cardinality to Actual Cardinality



```
SELECT /*+ GATHER_PLAN_STATISTICS */  
      COUNT(*)  
  FROM sales s, customers c  
 WHERE s.cust_id = c.cust_id  
   AND substr(c.cust_state_province,1,2)='CA';
```

```
SELECT * FROM table(dbms_xplan.display_cursor(format=>'ALLSTATS LAST'));
```

Id	Operation	Name	E-Rows	A-Rows	Buffer
0	SELECT STATEMENT		1	1	15984
1	SORT AGGREGATE		1	1	15984
2	NESTED LOOPS		1471	154K	15984
3	TABLE ACCESS FULL	CUSTOMER	100	6656	15984
* 4	BITMAP CONVERSION COUNT		4	154K	2914
* 5	INDEX RANGE SCAN	SALES_CUST_BIX			13034

Why Is The Cardinality Estimates So Wrong?

Start With Cardinality Estimate of the Table on the Left-Hand Side of Join

- Customers is the table on the left hand side of the join & predicates is

```
substr(c.cust_state_province,1,2) = 'CA'
```

Why Is The Cardinality Estimates So Wrong?

Start With Cardinality Estimate of the Table on the Left-Hand Side of Join

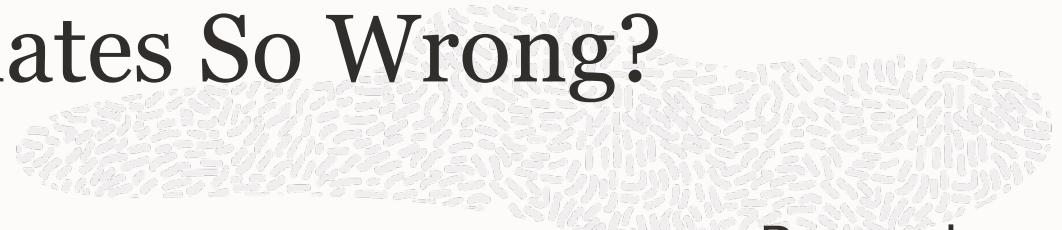
- Customers is the table on the left hand side of the join & predicates is

```
substr(c.cust_state_province,1,2) = 'CA' ← Function  
                                         wrapped column
```

- Optimizer doesn't know how function affects values in the column

Why Is The Cardinality Estimates So Wrong?

Function Wrapped Column Blinds the Optimizer



```
SELECT COUNT(*)  
FROM   customers  
WHERE  substr(c.cust_state_province,1,2) = 'CA';
```

Remember
CUSTOMERS table
has 10,000 rows

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			2 (100)
1	SORT AGGREGATE		1	
* 4	TABLE ACCESS FULL	CUSTOMER	100	2 (0)

- Optimizer doesn't know how function affects values in the column
- Optimizer guesses the cardinality to be 1% of rows

Tell the Optimizer How Function Affects Column Values

Create Extended Statistics



```
SELECT
dbms_stats.create_extended_stats(null,'CUSTOMERS',
        '(SUBSTR(CUST_STATE_PROVINCE,1,2))')
FROM dual;

DBMS_STATS.CREATE_EXTENDED_STATS(NULL, 'CUSTOMERS' , '(SUBSTR(CUST_STATE_PROVINCE,1,2))')
```

SYS_STUAJYEDA\$07W8CRW1A18K4Q_G



Tell the Optimizer How Function Affects Column Values

Identifying Extended Statistics

```
SELECT column_name, num_distinct, histogram  
FROM user_tab_col_statistics  
WHERE table_name='CUSTOMERS';
```

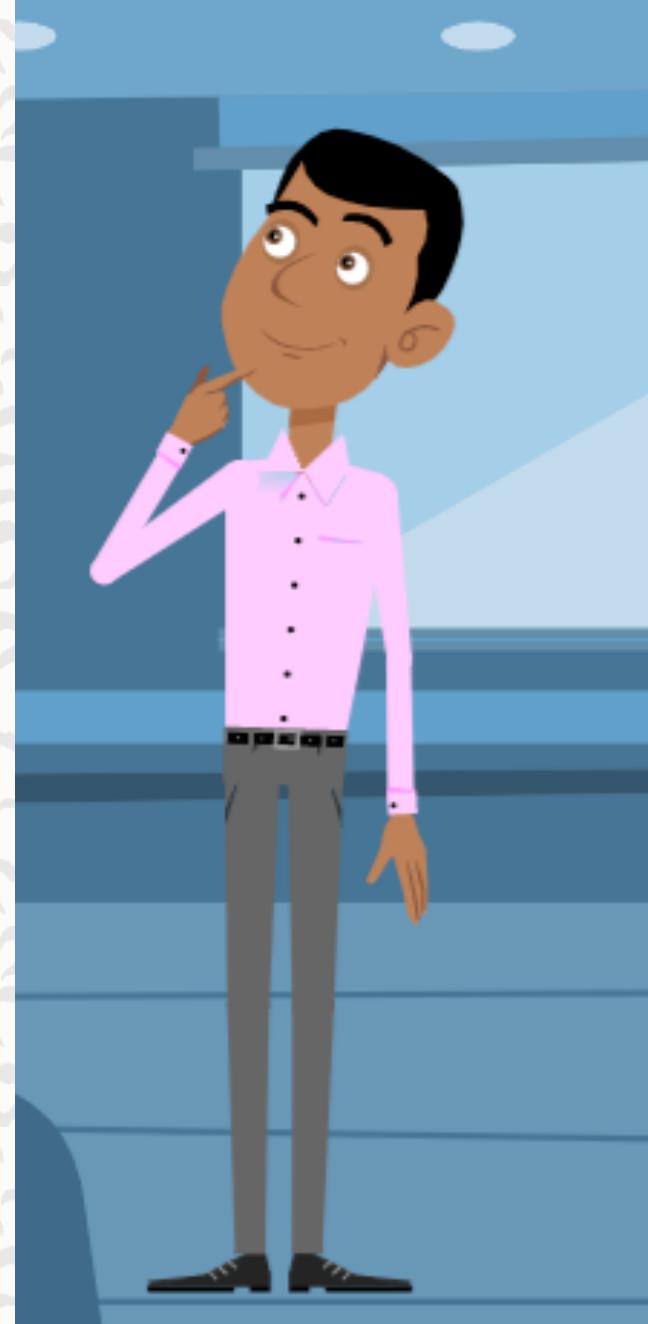
COLUMN_NAME	NUM_DISTINCT	HISTOGRAM
SYS_STUAJYEDA\$07W8CRW1A18K4Q_G	84	FREQUENCY
C_CUST_ID	335	HYBRID
CUST_YEAR_OF_BIRTH	69	FREQUENCY
CUST_VALID	2	FREQUENCY
CUST_TOTAL_ID	1	FREQUENCY
CUST_TOTAL	1	FREQUENCY
CUST_STREET_ADDRESS	340	HYBRID
:		

Fixing Cardinality Estimate on Left-Hand Side Automatic Produces a Hash Join Plan

```
SELECT COUNT(*)  
FROM   sales s, customers c  
WHERE  s.cust_id = c.cust_id  
AND    substr(c.cust_state_province,1,2)='CA';
```

Id	Operation	Name	Rows	Cost (%CPU)
0	SELECT STATEMENT			1114 (100)
1	SORT AGGREGATE		1	
2	HASH JOIN		14732	1114 (2)
* 4	TABLE ACCESS FULL	CUSTOMER	6656	406 (1)
* 5	TABLE ACCESS FULL	SALES	10M	703 (2)

Problem Statement 4



Expected Partition Pruning

```
SELECT    SUM(s.amount_sold)
FROM      sales s,
WHERE     TO_CHAR(s.time_id,'YYYYMMDD')
          BETWEEN '20190101' AND '20191231';
```

SALES table has 10 million rows

SALES is partitioned on the time_id column on a quarterly basis and 28 partitions

Expected Partition Pruning But Didn't Get It

```
SELECT    SUM(s.amount_sold)
FROM      sales s,
WHERE     TO_CHAR(s.time_id,'YYYYMMDD')
          BETWEEN '20190101' AND '20191231';
```

SALES table has 10 million rows

SALES is partitioned on the time_id column on a quarterly basis and 28 partitions

Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	PARTITION RANGE ALL		10M	552 (7)	1	28
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	1	28

Predicate Information (identified by operation id):

```
3 - filter((TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')>='20190101' AND
           TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')<='20191231'))
```

Expected Partition Pruning But Didn't Get It

```
SELECT      SUM(s.amount_sold)
FROM        sales s,
WHERE       TO_CHAR(s.time_id,'YYYYMMDD')
            BETWEEN '20190101' AND '20191231';
```

SALES table has 10 million rows

SALES is partitioned on the time_id column on a quarterly basis and 28 partitions

Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	PARTITION RANGE ALL		10M	552 (7)	1	28
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	1	28

WHY HAS AN ADDITIONAL INTERNAL_FUNCTION BEEN ADDED TO OUR PREDICATE?

3 - filter((TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')>='20190101' AND TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"),'YYYYMMDD')<='20191231'))



Expected Partition Pruning But Didn't Get It

Expected query to access only 4 partitions but its accesses all

- INTERNAL_FUNCTION typically means a data type conversion has occurred

Expected Partition Pruning But Didn't Get It

Expected query to access only 4 partitions but its accesses all

- INTERNAL_FUNCTION typically means a data type conversion has occurred
- Data type conversion needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`

Expected Partition Pruning But Didn't Get It

Expected query to access only 4 partitions but its accesses all

- INTERNAL_FUNCTION typically means a data type conversion has occurred
- Data type conversion needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`
- `TIME_ID` is actually a date column
- Optimizer has no idea how function will effect the values in `TIME_ID` column
- Optimizer can't determine which partitions will be accessed now

Expected Partition Pruning But Didn't Get It

Expected query to access only 4 partitions but its accesses all

- INTERNAL_FUNCTION typically means a data type conversion has occurred
- Data type conversion needed when column type & predicate type don't match
- Predicate is `TO_CHAR(s.time_id, 'YYYYMMDD')`
- `TIME_ID` is actually a date column
- Optimizer has no idea how function will effect the values in `TIME_ID` column
- Optimizer can't determine which partitions will be accessed now

SIMPLE FIX: USING INVERSE FUNCTION ON OTHER SIDE OF PREDICATE

Solution Use Inverse Function On Other Side Of Predicate



```
SELECT SUM(s.amount_sold)
FROM sales s,
WHERE s.time_id BETWEEN TO_DATE('20190101', 'YYMMDD') AND
TO_DATE(' 20191231', 'YYMMDD');
```

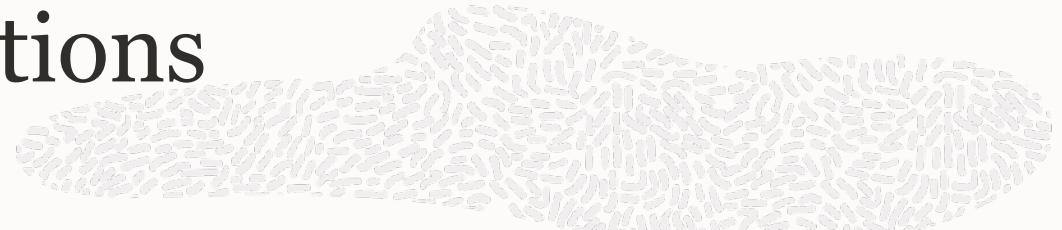
Id	Operation	Name	ROWS	Cost (%CPU)	Pstart	Pstop
0	SELECT STATEMENT			552 (100)		2
1	SORT AGGREGATE		13			2
2	PARTITION RANGE ALL		10M	552 (7)	9	12
* 3	TABLE ACCESS FULL	SALES	10M	552 (7)	9	12

Predicate Information (identified by operation id):

```
3 - Access("S"."TIME_ID">>='20190101' TO_DATE('20190101','YYMMDD') AND
TO_DATE(' 20191231','YYMMDD'))
```

Important Tip on Using Functions

Using inverse function on other side of predicate

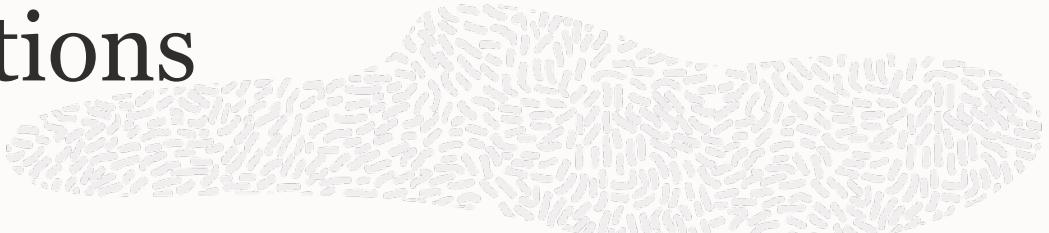


Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents index use or partition pruning

Important Tip on Using Functions

Using inverse function on other side of predicate

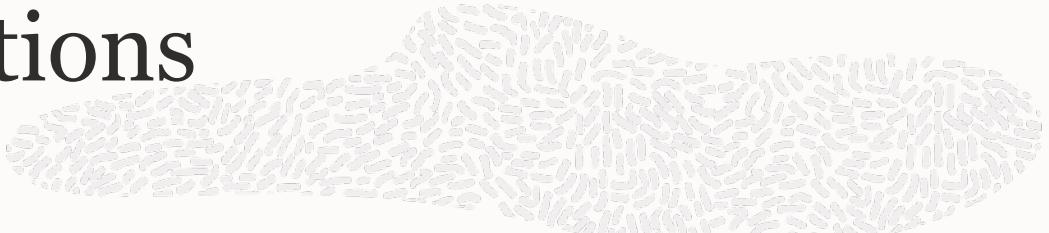


Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents index use or partition pruning
- For function-based index to be considered, use that exact function as specified in index

Important Tip on Using Functions

Using inverse function on other side of predicate



Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents index use or partition pruning
- For function-based index to be considered, use that exact function as specified in index
- If multiple predicates involve the same columns, write predicates such that they share common expressions For example,

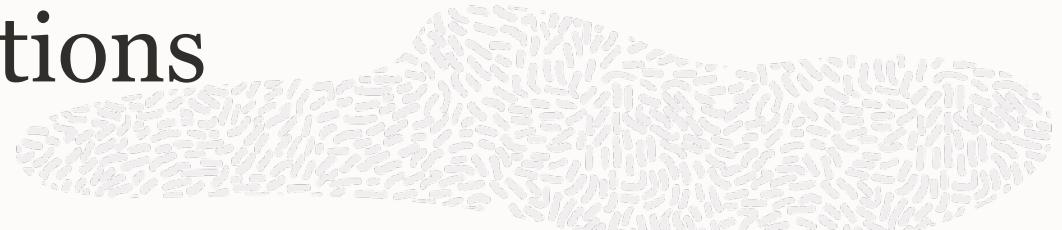
WHERE $f(a) = b$
AND $a = c$

Should be
rewritten as

WHERE $a = \text{inv_}f(b)$
AND $a = c$

Important Tip on Using Functions

Using inverse function on other side of predicate



Keep the following in mind when deciding where to place the function

- Try to place functions on top of constants (literals, binds) rather than on columns
- Avoid using a function on index columns or partition keys as it prevents index use or partition pruning
- For function-based index to be considered, use that exact function as specified in index
- If multiple predicates involve the same columns, write predicates such that they share common expressions For example,

WHERE $f(a) = b$
AND $a = c$

Should be
rewritten as

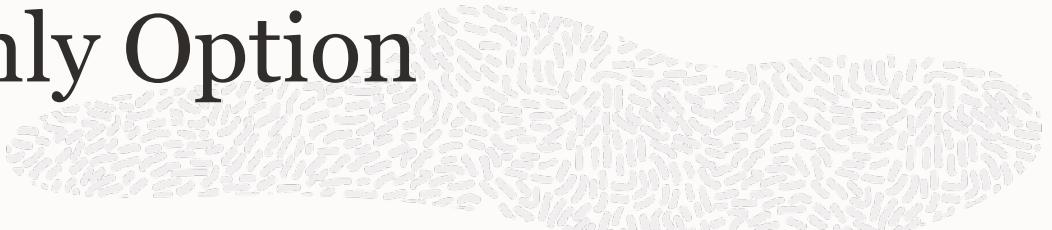
WHERE $a = \text{inv_}f(b)$
AND $a = c$

This will allow transitive predicate $c = \text{inv_}f(b)$ to be added by the optimizer

Quick Tips If You Must Use Hints



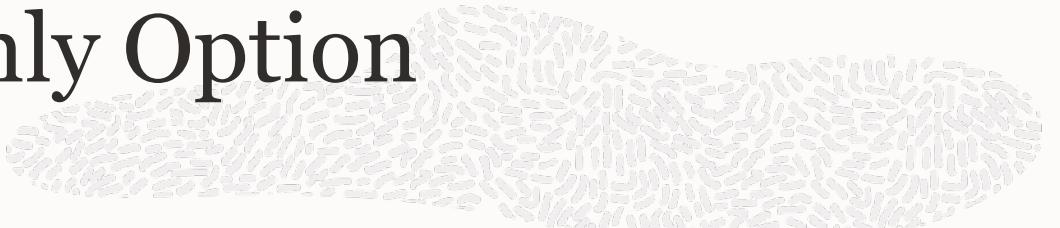
Some Times a Hint is Your Only Option



What are Optimizer Hints

- Hints allow you to influence the Optimizer when it has to choose between several possibilities
- A hint is a directive that will be followed when applicable

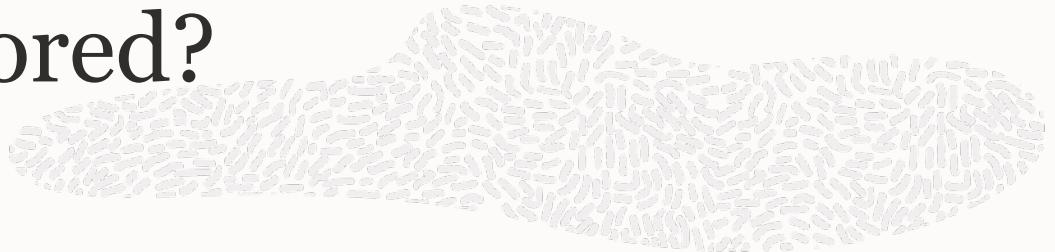
Some Times a Hint is Your Only Option



What are Optimizer Hints

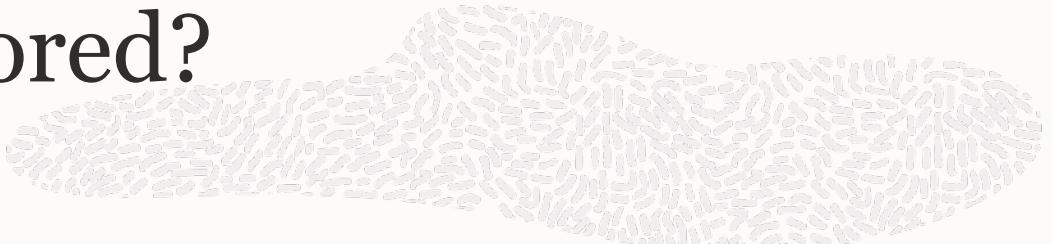
- Hints allow you to influence the Optimizer when it has to choose between several possibilities
- A hint is a directive that will be followed when applicable
- Can influence everything from the Optimizer mode used to each operation in the execution
- Automatically means the Cost Based Optimizer will be used
- Only exception is the RULE hint but it must be used alone

Why are Optimizer Hints Ignored?



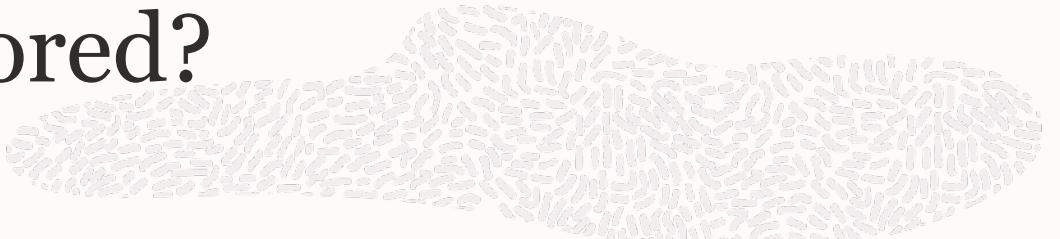
- An Optimizer hint is never ignored

Why are Optimizer Hints Ignored?



- An Optimizer hint is never ignored
- Hints only evaluated when they apply to a decision that has to be made
- Often times hint that's aren't used because they are irrelevant or not legal

Why are Optimizer Hints Ignored?



- An Optimizer hint is never ignored
- Hints only evaluated when they apply to a decision that has to be made
- Often times hint that's aren't used because they are irrelevant or not legal

In **10053 trace file** you will find:

“Dumping Hints

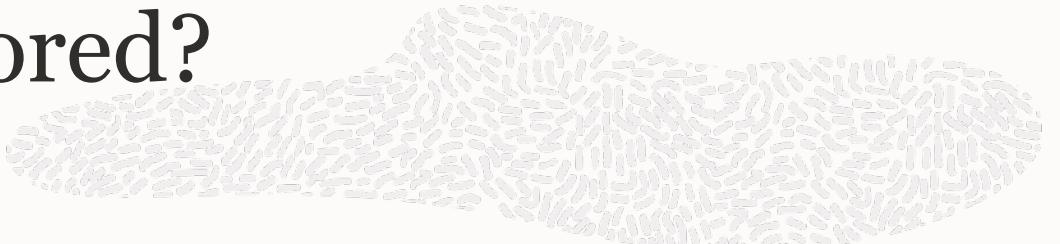
=====

Atom_hint=(@=0X124360178 err=0 resol=0 used=1 token=454 org=1 lvl=1 txt=ALL_ROWS)

Atom_hint=(@=0X2af785e0c260 err=0 resol=1 used=1 token=448 org=1 lvl=3 txt=FULL (“E”))

===== END SQL Statement Dump =====

Why are Optimizer Hints Ignored?



- An Optimizer hint is never ignored
- Hints only evaluated when they apply to a decision that has to be made
- Often times hint that's aren't used because they are irrelevant or not legal

In **10053 trace file** you will find:

“Dumping Hints

=====

Atom_hint=(@=0X124360178 err=0 resol=0 used=1 token=454 org=1 lvl=1 txt=ALL_ROWS)

Atom_hint=(@=0X2af785e0260 err=0 resol=1 used=1 token=448 org=1 lvl=3 txt=FULL (“E”))

===== == END SQL Statement =====

ERR indicates if there is
an error with hint

USED indicates the hint was used during the
evaluation of the part of the plan it pertains to But
Doesn’t mean the final plan will reflect it

Why are Optimizer Hints Ignored?

New Hint Usage Report in 19c

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				9 (100)	
1	HASH GROUP BY		1	52	9 (23)	00:00:01
* 2	HASH JOIN		1	52	8 (13)	00:00:01
3	PARTITION RANGE SINGLE		2	34	2 (0)	00:00:01
* 4	TABLE ACCESS FULL	SALES	2	34	2 (0)	00:00:01
* 5	TABLE ACCESS FULL	CUSTOMERS	13	455	5 (0)	00:00:01

Predicate Information (identified by operation id):

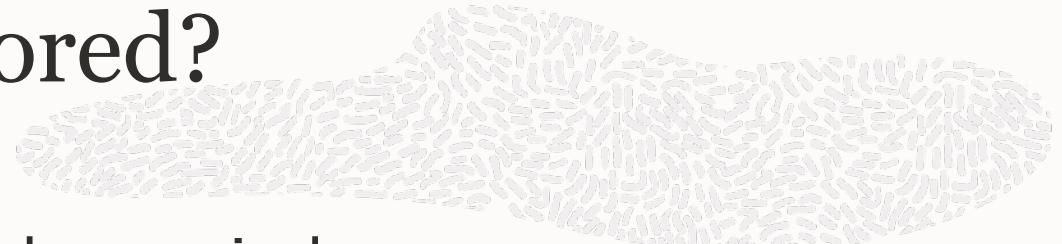
```
2 - access("C"."CUST_ID"="S"."CUST_ID")
4 - filter("S"."TIME_ID">'30-DEC-99')
5 - filter(("C"."CUST_CITY"='Los Angeles' AND "CUST_SEX"=1))
```

Hint Report (identified by operation id / Query Block Name / Object Alias):
Total hints for statement: 1 (U - Unused (1))

```
4 - SEL$1 / S@SEL$1
U - USE NL(s)
```

New hint info under the plan
in 19c with
DBMS_XPLAN.DISPLAY_CURSOR

Why are Optimizer Hints Ignored?



Statement 1

Employees table has a unique index called pk_emp index

```
SELECT /*+ INDEX(e emp_pk) */ * FROM employees e;
```

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	8743	(1)
1	TABLE ACCESS FULL	EMPLOYEES	2579	8743	(1)

Why are Optimizer hints ignored?

Syntax and Spelling

Employees table has a unique index called **pk_emp** index

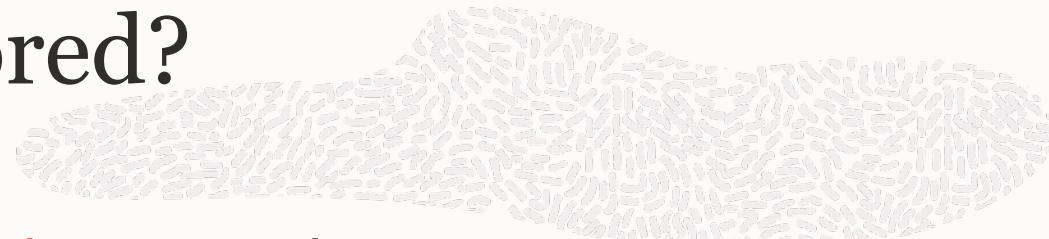
```
SELECT /*+ INDEX(e emp_pk) */ * FROM employees e;
```

Id	Operation	Name	Rows			
0	SELECT STATEMENT		2579	8743	(1)	
1	TABLE ACCESS FULL	EMPLOYEES	2579	8743	(1)	

Invalid hint the index name is wrong

Why are Optimizer hints ignored?

Syntax and Spelling



Employees table has a unique index called **pk_emp** index

```
SELECT /*+ INDEX(e emp_pk) */ * FROM employees e;
```

Id	Operation	Name	Rows	COST (%CPU)	
0	SELECT STATEMENT		2579	8743	(1)
1	TABLE ACCESS FULL	EMPLOYEES	2579	8743	(1)

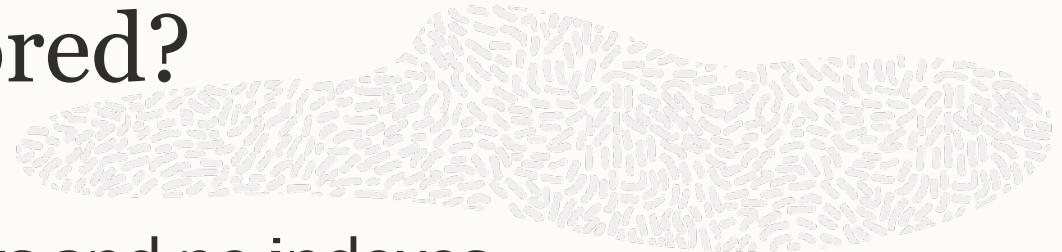
Hint Report (identified by operation id / Query Block Name / Object Alias):
Total hints for statement: 1 (U - Unused (1))

```
1 - SEL$1 /S@SEL$1
```

```
U - index (e emp_pk) / index specified in the hint doesn't exist
```

19^c

Why are Optimizer hints ignored?



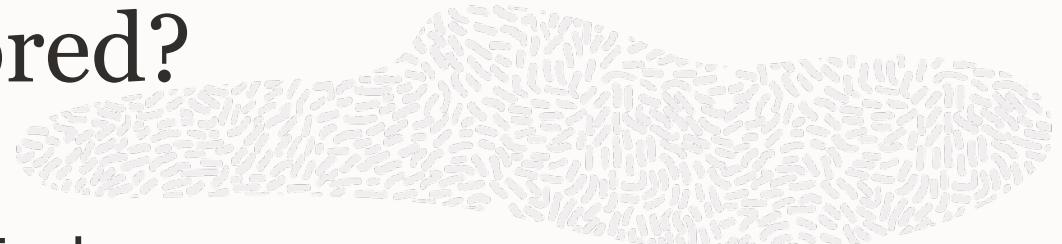
The MY_PROMOTIONS table has 5,000 rows and no indexes

```
SELECT /*+ INDEX(p) */ Count(*)
FROM   my_promotions p
WHERE  promo_category = 'TV'
AND    promo_begin_date = '05-OCT-17';
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				3 (100)	
1	SORT AGGREGATE		1	15		
*	TABLE ACCESS FULL	MY_PROMOTIONS	1	15	3 (0)	00:00:01

Why are Optimizer hints ignored?

Invalid hint



Specifying an index hint on a table with no indexes

```
SELECT /*+ INDEX(p) */ Count(*)
```

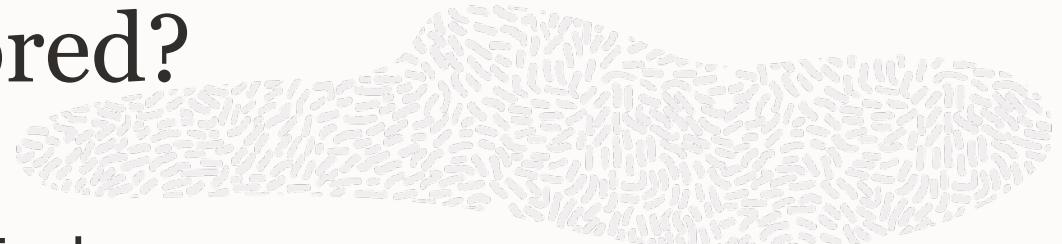
```
FROM my_promotions p
WHERE promo_category = 'TV'
AND promo_begin_date = '05-OCT-17';
```

Invalid hint because no indexes exist on the table

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				3 (100)	
1	SORT AGGREGATE		1	15		
* 2	TABLE ACCESS FULL	MY PROMOTIONS	1	15	3 (0)	00:00:01

Why are Optimizer hints ignored?

Invalid hint



Specifying an index hint on a table with no indexes

```
SELECT /*+ INDEX(p) */ Count(*)
```

```
FROM my_promotions p
WHERE promo_category = 'TV'
AND promo_begin_date = '05-OCT-17';
```

Invalid hint because no indexes exist on the table

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				3 (100)	
1	SORT AGGREGATE		1	15		
* 2	TABLE ACCESS FULL	MY PROMOTIONS	1	15	3 (0)	00:00:01

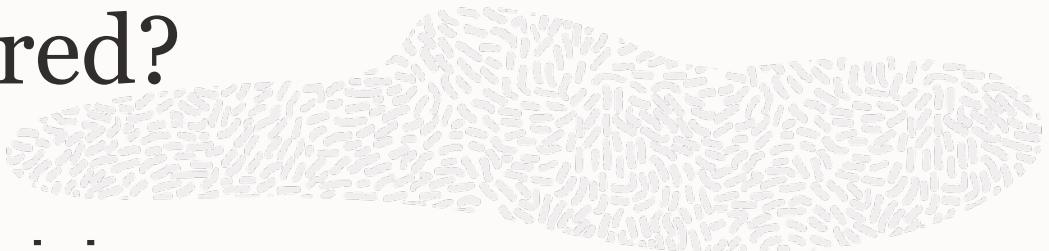
Hint Report (identified by operation id / Query Block Name / Object Alias):
Total hints for statement: 1 (U - Unused (1))

```
2 - SEL$1 /P@SEL$1
      U - INDEX(p)
```

19c



Why are Optimizer hints ignored?



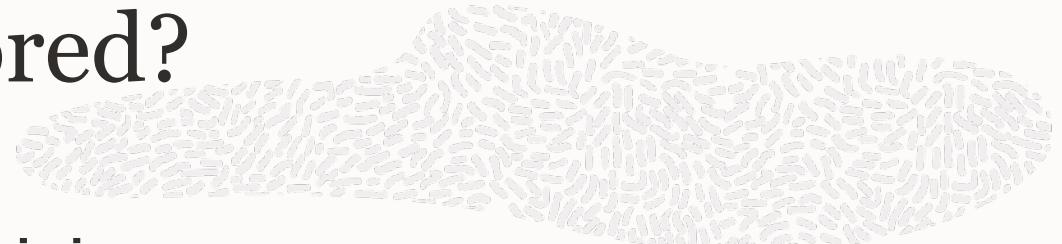
Requesting a hash join hint for non-equality join

```
SELECT /*+ USE_HASH(e s) */ e.first_name, e.last_name
FROM   employees e, salary_grade s
WHERE  e.salary BETWEEN s.low_sal AND s.high_sal;
```

Id Operation	Name	Rows	Bytes	Cost (CPU)
0 SELECT STATEMENT		1	45	4 (0)
1 NESTED LOOPS		1	45	4 (0)
/* 2 TABLE ACCESS STORAGE FULL	SALARY_GRADE	1	26	2 (0)
/* 3 TABLE ACCESS STORAGE FULL	EMPLOYEES	1	19	2 (0)

Why are Optimizer hints ignored?

Illegal hint



Specifying a hash join hint for non-equality join

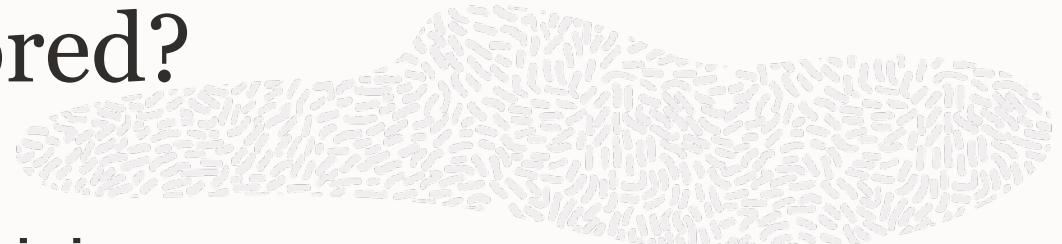
```
SELECT /*+ USE_HASH(e s) */ e.first_name, e.last_name  
FROM   employees e, salary_grade s  
WHERE  e.salary BETWEEN s.low_sal AND s.high_sal;
```

Id Operation	Name	Rows	Bytes	Cost (CPU)
0 SELECT STATEMENT		1	45	4 (0)
1 NESTED LOOPS		1	45	4 (0)
* 2 TABLE ACCESS STORAGE FULL	SALARY_GRADE	1	26	2 (0)
* 3 TABLE ACCESS STORAGE FULL	EMPLOYEES	1	19	2 (0)

Illegal hint because a hash join can't be used for a non-equality join predicate

Why are Optimizer hints ignored?

Illegal hint



Specifying a hash join hint for non-equality join

```
SELECT /*+ USE_HASH(e s) */ e.first_name, e.last_name  
FROM   employees e, salary_grade s  
WHERE  e.salary BETWEEN s.low_sal AND s.high_sal;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		1	45	4 (0)
1	NESTED LOOPS		1	45	4 (0)
* 2	TABLE ACCESS STORAGE FULL	SALARY_GRADE	1	26	2 (0)
* 3	TABLE ACCESS STORAGE FULL	EMPLOYEES	1	19	2 (0)

Illegal hint because a hash join can't be used for a non-equality join predicate

Hint Report (identified by operation id / Query Block Name / Object Alias):
Total hints for statement: 1 (U - Unused (1))

```
2 - SEL$1 /P@SEL$1  
U - USE_HASH(e s)
```

19c

Tips To Remember



Statistics

Always check and correct
your cardinality estimates
first

Tips To Remember



Statistics

Always check and correct
your cardinality estimates
first



SQL Statement

Look at how the SQL
statement is written and how
the predicates are being used

Tips To Remember

1

Statistics

Always check and correct
your cardinality estimates
first

2

SQL Statement

Look at how the SQL
statement is written and how
the predicates are being used

3

Hints

Only add hints as a last
resort and confirm your hint
is really being used

Helpful Links

 <https://sqlmaria.com>

 <https://twitter.com/@SQLMaria>

 <https://www.facebook.com/SQLMaria>



ORACLE